

Data Science for Causal Inference

Ryan T. Moore

American University

The Lab @ DC

2024-07-15

Table of contents I

Introductions

Data Science in Causal Inference

Heterogeneous Treatment Effects

Variable Selection

Introductions

About Me

- ▶ Associate Prof of Government
(American University)
- ▶ Associate Director, Center for Data Science
(American University)
- ▶ Senior Social Scientist
(The Lab @ DC)
- ▶ Fellow in Methodology
(US Office of Evaluation Sciences: “OES”)

About Me

- ▶ Associate Prof of Government
(American University)
- ▶ Associate Director, Center for Data Science
(American University)
- ▶ Senior Social Scientist
(The Lab @ DC)
- ▶ Fellow in Methodology
(US Office of Evaluation Sciences: “OES”)
- ▶ Research agenda: political methodology,
causal inference, experimental design,
experiments in public policy

About You!

► Name?

About You!

▶ Name?

▶ Role?

About You!

- ▶ Name?
- ▶ Role?
- ▶ Interests?

About You!

- ▶ Name?
- ▶ Role?
- ▶ Interests?
- ▶ Olympic sport you look forward to?

Plan

- ▶ Data Science in Causal Inference

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders
- ▶ Modern difference-in-difference designs

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders
- ▶ Modern difference-in-difference designs
 - ▶ Canonical DiD

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders
- ▶ Modern difference-in-difference designs
 - ▶ Canonical DiD
 - ▶ Multiple time periods

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders
- ▶ Modern difference-in-difference designs
 - ▶ Canonical DiD
 - ▶ Multiple time periods
 - ▶ Staggered adoption

Plan

- ▶ Data Science in Causal Inference
 - ▶ Models
 - ▶ Heterogeneous treatment effects
 - ▶ Variable selection
- ▶ Sensitivity
 - ▶ Model specification
 - ▶ Unobservable parameter
 - ▶ Unobserved confounders
- ▶ Modern difference-in-difference designs
 - ▶ Canonical DiD
 - ▶ Multiple time periods
 - ▶ Staggered adoption
 - ▶ Calloway-Sant'Anna approach

Data Science in Causal Inference

Causal Inference Approaches

The “potential outcomes” framework:

Causal Inference Approaches

The “potential outcomes” framework:

Citizen	Canvass?	Would Enroll if Canvass?	Would Enroll if No Canvass?	Enroll
1	Yes	Yes		Yes
2	Yes			Yes
3	No			No
4	No			No

Causal Inference Approaches

The “potential outcomes” framework:

Citizen	Canvass?	Would Enroll if Canvass?	Would Enroll if No Canvass?	Enroll
1	Yes	Yes		Yes
2	Yes	Yes		Yes
3	No			No
4	No			No

Causal Inference Approaches

The “potential outcomes” framework:

Citizen	Canvass?	Would Enroll if Canvass?	Would Enroll if No Canvass?	Enroll
1	Yes	Yes		Yes
2	Yes	Yes		Yes
3	No		No	No
4	No			No

Causal Inference Approaches

The “potential outcomes” framework:

Citizen	Canvass?	Would Enroll if	Would Enroll if	Enroll
		Canvass?	No Canvass?	
1	Yes	Yes		Yes
2	Yes	Yes		Yes
3	No		No	No
4	No		No	No

Causal Inference Approaches

The “potential outcomes” framework:

Citizen	Canvass?	Would Enroll if	Would Enroll if	Enroll
		Canvass?	No Canvass?	
1	Yes	Yes	(Yes)	Yes
2	Yes	Yes	(No)	Yes
3	No	(Yes)	No	No
4	No	(No)	No	No

Causal Inference Approaches

The “potential outcomes” framework, more abstractly:

Unit i	Treatment T	$Y(1)$	$Y(0)$	Y^{obs}	True τ $Y(1) - Y(0)$
1	1	10		10	
2	1	20		20	
3	0		15	15	
4	0		5	5	

Causal Inference Approaches

The “potential outcomes” framework, more abstractly:

Unit i	Treatment T	$Y(1)$	$Y(0)$	Y^{obs}	True τ
					$Y(1) - Y(0)$
1	1	10	(10)	10	0
2	1	20	(10)	20	10
3	0	(40)	15	15	25
4	0	(20)	5	5	15

Causal Inference Approaches

The “potential outcomes” framework, more abstractly:

Unit i	Treatment T			Y^{obs}	True τ
		$Y(1)$	$Y(0)$		$Y(1) - Y(0)$
1	1	10	(10)	10	0
2	1	20	(10)	20	10
3	0	(40)	15	15	25
4	0	(20)	5	5	15
ATE = $\bar{\tau}$ =					$\frac{50}{4} = 12.5$

Causal Inference Approaches

The “potential outcomes” framework, more abstractly:

Unit i	Treatment T			Y^{obs}	True τ
		$Y(1)$	$Y(0)$		$Y(1) - Y(0)$
1	1	10	(10)	10	0
2	1	20	(10)	20	10
3	0	(40)	15	15	25
4	0	(20)	5	5	15
				$ATE = \bar{\tau} =$	$\frac{50}{4} = 12.5$
				$\widehat{ATE} = \hat{\tau} =$	$15 - 10 = 5$

Causal Inference Approaches

The “potential outcomes” framework, notation:

- ▶ Units indexed by i
- ▶ Treatment T_i or D_i or Z_i
- ▶ Outcome if treated $Y_i(1)$
- ▶ Outcome if control $Y_i(0)$
- ▶ True treatment effect $\tau_i = Y_i(1) - Y_i(0)$
- ▶ True average treatment effect
$$\bar{\tau} = \frac{1}{n} \sum_{i=1}^n (Y_i(1) - Y_i(0))$$
- ▶ Pre-treatment covariates \mathbf{X}

Causal Inference Approaches

The “potential outcomes” framework, notation:

- ▶ Units indexed by i
- ▶ Treatment T_i or D_i or Z_i
- ▶ Outcome if treated $Y_i(1)$
- ▶ Outcome if control $Y_i(0)$
- ▶ True treatment effect $\tau_i = Y_i(1) - Y_i(0)$
- ▶ True average treatment effect
$$\bar{\tau} = \frac{1}{n} \sum_{i=1}^n (Y_i(1) - Y_i(0))$$
- ▶ Pre-treatment covariates \mathbf{X}

(and we'll draw some DAG's, too)

Data Science Approaches

Three tasks of data science:

- ▶ Description

Data Science Approaches

Three tasks of data science:

- ▶ Description
- ▶ Prediction

Data Science Approaches

Three tasks of data science:

- ▶ Description
- ▶ Prediction
- ▶ Causal Inference

Data Science Approaches

Three tasks of data science:

- ▶ Description
- ▶ Prediction
- ▶ Causal Inference

Data Science Approaches

Three tasks of data science:

- ▶ Description
- ▶ Prediction
- ▶ Causal Inference

Models/algorithms central to all three.

Data Science Approaches

Three tasks of data science:

- ▶ Description
- ▶ Prediction
- ▶ Causal Inference

Models/algorithms central to all three.

Hernán, Hsu, and Healy (2019)

Data Science Approaches

Description

- ▶ Identifying patterns, etc.

Data Science Approaches

Description

- ▶ Identifying patterns, etc.
- ▶ E.g., clustering to discover groups

Data Science Approaches

Prediction

► Components

Data Science Approaches

Prediction

- ▶ Components
 - ▶ Inputs/outputs (predictors/outcomes, features/responses, ...)

Data Science Approaches

Prediction

- ▶ Components
 - ▶ Inputs/outputs (predictors/outcomes, features/responses, ...)
 - ▶ Mapping from inputs to outputs (linear model, decision tree, ...)

Data Science Approaches

Prediction

- ▶ Components
 - ▶ Inputs/outputs (predictors/outcomes, features/responses, ...)
 - ▶ Mapping from inputs to outputs (linear model, decision tree, ...)
 - ▶ Metric for evaluating mapping

Data Science Approaches

Prediction

- ▶ Components
 - ▶ Inputs/outputs (predictors/outcomes, features/responses, ...)
 - ▶ Mapping from inputs to outputs (linear model, decision tree, ...)
 - ▶ Metric for evaluating mapping
- ▶ With these, model machine learning does the work

Data Science Approaches

Prediction

- ▶ Components
 - ▶ Inputs/outputs (predictors/outcomes, features/responses, ...)
 - ▶ Mapping from inputs to outputs (linear model, decision tree, ...)
 - ▶ Metric for evaluating mapping
- ▶ With these, model machine learning does the work
- ▶ E.g., regression, random forests, neural networks, ...

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure
 - ▶ Not all inputs treated same

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure
 - ▶ Not all inputs treated same
 - ▶ T v. \mathbf{X} – very different!

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure
 - ▶ Not all inputs treated same
 - ▶ T v. \mathbf{X} – very different!
 - ▶ (the more knowledge, the better!)

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure
 - ▶ Not all inputs treated same
 - ▶ T v. \mathbf{X} – very different!
 - ▶ (the more knowledge, the better!)
 - ▶ (alternative: solve fundamental problem of causal inference!)

Data Science Approaches

Causal Inference

- ▶ Potential outcomes/counterfactual/interventionist perspective
- ▶ Requires *expertise* different to description/prediction
- ▶ Requires more than summary statistics, metrics, etc.
- ▶ Requires some knowledge of causal structure
 - ▶ Not all inputs treated same
 - ▶ T v. \mathbf{X} – very different!
 - ▶ (the more knowledge, the better!)
 - ▶ (alternative: solve fundamental problem of causal inference!)
- ▶ E.g., experiments, observational causal designs, ...

Causal Inference with Machine Learning

Causal Inference with Machine Learning



Jake M. Grumbach

@JakeMGrumbach

...

I finally found it in real life: the consultant who runs OLS in Excel and calls it machine learning

9:17 AM · Jan 31, 2019 · Twitter for iPhone

54 Retweets **7** Quote Tweets **511** Likes



Causal Inference with Machine Learning



Jake M. Grumbach

@JakeMGrumbach



I finally found it in real life: the consultant who runs OLS in Excel and calls it machine learning

9:17 AM · Jan 31, 2019 · Twitter for iPhone

54 Retweets **7** Quote Tweets **511** Likes



(OK, not “machine learning”, perhaps, but *models* at least ...)

Causal Inference with Models

Loaded two datasets:



```
str(df1)
```

```
tibble [100 x 3] (S3: tbl_df/tbl/data.frame)
 $ covariate: num [1:100] -0.622 1.137 -0.238 1.529 -0.154
 $ exposure : num [1:100] 0.0332 0.3627 0.2422 1.4633 0.779
 $ outcome  : num [1:100] -0.429 2.675 -0.647 2.238 1.044
```

```
str(df2)
```

```
tibble [100 x 3] (S3: tbl_df/tbl/data.frame)
 $ exposure : num [1:100] 0.4862 0.0653 -1.4021 -0.546 -0.4
 $ outcome  : num [1:100] 1.706 0.669 -1.597 -1.733 0.617
 $ covariate: num [1:100] 2.24 0.924 -0.999 -2.343 0.207
```

Causal Inference with Models



Causal Inference with Models

Model each

```
lm_df1 <- lm(outcome ~ exposure, data = df1)
lm_df2 <- lm(outcome ~ exposure, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	(Intercept)	-0.00671	0.120
2	df1	exposure	0.996	0.0927
3	df2	(Intercept)	0.133	0.0890
4	df2	exposure	1.00	0.0841

Causal Inference with Models

Model each

```
lm_df1 <- lm(outcome ~ exposure, data = df1)
lm_df2 <- lm(outcome ~ exposure, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	(Intercept)	-0.00671	0.120
2	df1	exposure	0.996	0.0927
3	df2	(Intercept)	0.133	0.0890
4	df2	exposure	1.00	0.0841

► Both cases: effect of exposure ≈ 1 .

Causal Inference with Models

Model each

```
lm_df1 <- lm(outcome ~ exposure, data = df1)
lm_df2 <- lm(outcome ~ exposure, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	(Intercept)	-0.00671	0.120
2	df1	exposure	0.996	0.0927
3	df2	(Intercept)	0.133	0.0890
4	df2	exposure	1.00	0.0841

- ▶ Both cases: effect of exposure ≈ 1 .
- ▶ Is this good?

Causal Inference with Models

Model each

```
lm_df1 <- lm(outcome ~ exposure, data = df1)
lm_df2 <- lm(outcome ~ exposure, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	(Intercept)	-0.00671	0.120
2	df1	exposure	0.996	0.0927
3	df2	(Intercept)	0.133	0.0890
4	df2	exposure	1.00	0.0841

- ▶ Both cases: effect of exposure ≈ 1 .
- ▶ Is this good?
- ▶ What if we adjust for covariate?

Causal Inference with Models

```
lm_df1_adj <- lm(outcome ~ exposure + covariate, data = df1)
lm_df2_adj <- lm(outcome ~ exposure + covariate, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	exposure	0.501	0.108
2	df1	covariate	0.970	0.147
3	df2	exposure	0.554	0.0990
4	df2	covariate	0.385	0.0598

► Both cases: effect of exposure ≈ 0.5 .

Causal Inference with Models

```
lm_df1_adj <- lm(outcome ~ exposure + covariate, data = df1)
lm_df2_adj <- lm(outcome ~ exposure + covariate, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	exposure	0.501	0.108
2	df1	covariate	0.970	0.147
3	df2	exposure	0.554	0.0990
4	df2	covariate	0.385	0.0598

- ▶ Both cases: effect of exposure ≈ 0.5 .
- ▶ Is this good?

Causal Inference with Models

```
lm_df1_adj <- lm(outcome ~ exposure + covariate, data = df1)
lm_df2_adj <- lm(outcome ~ exposure + covariate, data = df2)
```

A tibble: 4 x 4

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	exposure	0.501	0.108
2	df1	covariate	0.970	0.147
3	df2	exposure	0.554	0.0990
4	df2	covariate	0.385	0.0598

- ▶ Both cases: effect of exposure ≈ 0.5 .
- ▶ Is this good?
- ▶ Which is correct? $\beta = 1$? $\beta = 0.5$?

Causal Inference with Models

```
lm_df1_adj <- lm(outcome ~ exposure + covariate, data = df1)
lm_df2_adj <- lm(outcome ~ exposure + covariate, data = df2)
```

```
# A tibble: 4 x 4
```

	data	term	estimate	std.error
	<chr>	<chr>	<dbl>	<dbl>
1	df1	exposure	0.501	0.108
2	df1	covariate	0.970	0.147
3	df2	exposure	0.554	0.0990
4	df2	covariate	0.385	0.0598

- ▶ Both cases: effect of exposure ≈ 0.5 .
- ▶ Is this good?
- ▶ Which is correct? $\beta = 1$? $\beta = 0.5$?
- ▶ *Should* we adjust for covariate?

Causal Inference with Models

There is nothing in the data that tells us.



Causal Inference with Models

There is nothing in the data that tells us. ☹

Causal Inference with Models

There is nothing in the data that tells us. ☹

Here are the true structures:



Causal Inference with Models

When know structures, adjustment sets for unbiasedness differ:

- ▶ df1: confounding \Rightarrow **adjust for X**
- ▶ df2: collider \Rightarrow **do not adjust for X**

```
g_conf <- dagitty("dag{ x -> y ; x <- c -> y }")  
g_coll <- dagitty("dag{ x -> y ; x -> c <- y }")
```

```
adjustmentSets(g_conf, "x", "y")
```

```
{ c }
```



```
adjustmentSets(g_coll, "x", "y")
```

```
{ }
```

Causal Inference with Models

When know structures, adjustment sets for unbiasedness differ:



- ▶ df1: confounding \Rightarrow **adjust for X**
- ▶ df2: collider \Rightarrow **do not adjust for X**



Causal Inference with Models

When know structures, adjustment sets for unbiasedness differ:

- ▶ df1: confounding \Rightarrow **adjust for X**
- ▶ df2: collider \Rightarrow **do not adjust for X**



(Data from D'Agostino McGowan (2023))

Causal Inference with Models

- ▶ Importance of identifying “pre-treatment covariates”, “proper covariates”; doing “design before analysis”

Causal Inference with Models

- ▶ Importance of identifying “pre-treatment covariates”, “proper covariates”; doing “design before analysis”
- ▶ Importance of experiments: strong knowledge about (part of) causal structure

Causal Inference with Models

- ▶ Importance of identifying “pre-treatment covariates”, “proper covariates”; doing “design before analysis”
- ▶ Importance of experiments: strong knowledge about (part of) causal structure
- ▶ Causal inference is critical to scientific questions, and separate from prediction

Causal Inference with Models

- ▶ Importance of identifying “pre-treatment covariates”, “proper covariates”; doing “design before analysis”
- ▶ Importance of experiments: strong knowledge about (part of) causal structure
- ▶ Causal inference is critical to scientific questions, and separate from prediction
- ▶ Though, methods from prediction can aid causal inference

Causal Inference with Models

- ▶ Importance of identifying “pre-treatment covariates”, “proper covariates”; doing “design before analysis”
- ▶ Importance of experiments: strong knowledge about (part of) causal structure
- ▶ Causal inference is critical to scientific questions, and separate from prediction
- ▶ Though, methods from prediction can aid causal inference
- ▶ (A perspective on “causal euphemisms”: Hernán (2018))

Approaches of Prediction and Causal Inference

Two Cultures, (Breiman 2001)

- ▶ *Data Models*: our “social science modeling”
- ▶ *Algorithmic Models*: our “data science algorithms”

Methods for Prediction and Causal Inference

- ▶ Cross-validation
- ▶ Regression/Decision trees
- ▶ Random forests

James et al. (2021)

Cross-validation

k -fold cross-validation

- ▶ Randomly partition data into k groups
- ▶ Apply method to $k - 1$ groups
- ▶ Use result to predict for left-out group
- ▶ Calculate $\text{MSE}_i = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- ▶ Calculate test error as average of the k MSE's:

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$

- ▶ Select model that minimises $CV_{(k)}$

CV for Linear Model

```
library(tidyverse)

## Make data

mk_data <- function(n = 90, n_folds = 10){

  df <- tibble(
    x1 = rnorm(n),
    x2 = rnorm(n),
    x3 = rnorm(n),
    y = 0.1 * x1 + 0.2 * x2 + 0.5 * x3 + rnorm(n),
    cv_fold = sample(rep(1:n_folds, (n / n_folds)))
  )

}

df <- mk_data()
```

CV for Linear Model

```
head(df)
```

```
# A tibble: 6 x 5
```

	x1	x2	x3	y	cv_fold
	<dbl>	<dbl>	<dbl>	<dbl>	<int>
1	-0.403	1.54	1.48	1.66	7
2	-0.164	-2.71	-0.897	-2.15	1
3	-0.594	-0.398	-0.591	0.147	4
4	1.33	0.129	-1.75	-0.793	8
5	-0.249	-0.544	1.40	-0.507	8
6	2.02	0.0630	1.05	-0.130	2

CV for Linear Model

```
head(df)
```

```
# A tibble: 6 x 5
```

	x1	x2	x3	y	cv_fold
	<dbl>	<dbl>	<dbl>	<dbl>	<int>
1	-0.403	1.54	1.48	1.66	7
2	-0.164	-2.71	-0.897	-2.15	1
3	-0.594	-0.398	-0.591	0.147	4
4	1.33	0.129	-1.75	-0.793	8
5	-0.249	-0.544	1.40	-0.507	8
6	2.02	0.0630	1.05	-0.130	2

```
table(df$cv_fold)
```

[illegible]

CV for Linear Model

```
cv_lm <- function(data, fmla){  
  
  n_folds <- max(data$cv_fold)  
  store_mses <- vector("numeric", length = n_folds)  
  
  for(idx in 1:n_folds){  
  
    df_train <- data |> filter(cv_fold != idx)  
    df_test <- data |> filter(cv_fold == idx)  
  
    lm_out <- lm(fmla, data = df_train)  
  
    predictions <- predict(lm_out, newdata = df_test)  
  
    store_mses[idx] <- mean((df_test$y - predictions)^2)}  
  
  test_error_cv_k <- mean(store_mses)  
  return(test_error_cv_k)
```

CV for Linear Model

```
cv_lm(data = df, fmla = y ~ x1 + x2)
```

```
[1] 1.468833
```

CV for Linear Model

```
cv_lm(data = df, fmla = y ~ x1 + x2)
```

```
[1] 1.468833
```

```
df <- mk_data()  
cv_lm(df, y ~ x1 + x2)
```

```
[1] 1.344094
```

CV for Linear Model

```
cv_lm(data = df, fmla = y ~ x1 + x2)
```

```
[1] 1.468833
```

```
df <- mk_data()  
cv_lm(df, y ~ x1 + x2)
```

```
[1] 1.344094
```

```
df <- mk_data()  
cv_lm(df, y ~ x1 + x2 + x3)
```

```
[1] 0.9552452
```

CV for Linear Model




Figure 1: MSE always less (better) for 3-variable model.

Regression Trees

- ▶ Partition predictor space into regions R_1, R_2, \dots, R_J .
- ▶ If unit falls in region R_j , use average outcome in R_j as predicted value: \hat{y}_{R_j}
- ▶ (For *decision* on discrete outcome, count votes in R_j)
- ▶ Goal: minimise residual sum of squares (RSS), just like LS regression:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

Regression Trees

How to define regions R_j ?

Regression Trees

How to define regions R_j ?

- ▶ Top-down, greedy recursive binary split
- ▶ At each step, find predictor and cut-point that minimise

$$\sum_{i:x \in R_1(j,s)} (y_i - \hat{y}_{R_1(j,s)})^2 + \sum_{i:x \in R_2(j,s)} (y_i - \hat{y}_{R_2(j,s)})^2$$

Regression Trees

- ▶ Overfitting is a potential problem

Regression Trees

- ▶ Overfitting is a potential problem
- ▶ Can we increase predictive quality by only using *part* of a tree?

Regression Trees

- ▶ Overfitting is a potential problem
- ▶ Can we increase predictive quality by only using *part* of a tree?
- ▶ “Pruning”

Regression Trees

Pruning

- ▶ Build a large tree

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T
- ▶ m : terminal node index

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T
- ▶ m : terminal node index
- ▶ Find subtree that minimises

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T
- ▶ m : terminal node index
- ▶ Find subtree that minimises

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T
- ▶ m : terminal node index
- ▶ Find subtree that minimises

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} \left(y_i - \hat{y}_{R_m} \right)^2 + \alpha |T|$$

Regression Trees

Pruning

- ▶ Build a large tree
- ▶ Select the subtree that gives least prediction error (via cross-validation)
- ▶ But, many possible subtrees, so penalise larger trees via α
- ▶ α : penalty parameter
- ▶ $|T|$: count of terminal nodes of T
- ▶ m : terminal node index
- ▶ Find subtree that minimises

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} \left(y_i - \hat{y}_{R_m} \right)^2 + \alpha |T|$$

Sum squared pred. error (plus penalty that grows with tree size) across units in region, then regions.

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then
 - 3a. Do 1 and 2 on all but k th fold

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then
 - 3a. Do 1 and 2 on all but k th fold
 - 3b. Predict for k th fold, calculate MSE for several values of α

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then
 - 3a. Do 1 and 2 on all but k th fold
 - 3b. Predict for k th fold, calculate MSE for several values of α
 - 3c. Get avg MSE for each α

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then
 - 3a. Do 1 and 2 on all but k th fold
 - 3b. Predict for k th fold, calculate MSE for several values of α
 - 3c. Get avg MSE for each α
 - 3d. Pick α to minimise MSE

Regression Trees

But, how to choose α ? (Use cross-validation.)

1. Build big tree on training data (with some minimum terminal node size)
2. For several values of α , find best subtree.
3. Find best value of α via CV. Create K folds. Then
 - 3a. Do 1 and 2 on all but k th fold
 - 3b. Predict for k th fold, calculate MSE for several values of α
 - 3c. Get avg MSE for each α
 - 3d. Pick α to minimise MSE
4. Using that α , select best subtree from Step 2

Example: Regression Tree

```
library(qss)
library(rsample)
library(tree)

data("MPs")
mps <- MPs |> mutate(age = yod - yob,
                     is_labour = if_else(party == "labour", 1, 0),
                     is_london = if_else(region == "Greater London", 1, 0),
                     is_winner = if_else(margin > 0, 1, 0))
select(ln.net, age, is_labour, is_london, is_winner) |>
na.omit()


set.seed(765076184)

mp_split <- initial_split(mps, prop = 0.7)

mp_train <- training(mp_split)
mp_test <- testing(mp_split)
```

Example: Regression Tree

```
plot(tree_mp)  
text(tree_mp)
```



Example: Regression Tree

Would pruning help?

```
cv_mps <- cv.tree(tree_mp, K = 10)  
  
plot(cv_mps$size, cv_mps$dev, type = "b")
```




Figure 3: Subtree size 2 minimises SSR

Example: Regression Tree

Would pruning help?

```
cv_mps <- cv.tree(tree_mp, K = 10)  
  
plot(cv_mps$size, cv_mps$dev, type = "b")
```




Figure 3: Subtree size 2 minimises SSR

Example: Regression Tree

```
prune_mps <- prune.tree(tree_mp, best = 2)  
  
plot(prune_mps)  
text(prune_mps)
```




Figure 4: The pruned tree

Example: Regression Tree

Predict for test set:

- ▶ MSE for pruned: 1.922
- ▶ MSE for full: 1.945

Example: Regression Tree

Predict for test set:

- ▶ MSE for pruned: 1.922
- ▶ MSE for full: 1.945

So, pruning helped us avoid some overfitting.

Example: Regression Tree

Predict for test set:

- ▶ MSE for pruned: 1.922
- ▶ MSE for full: 1.945

So, pruning helped us avoid some overfitting.

(Typical pred error of $\sqrt{1.922} \approx 1.386$)

Example: Regression Tree

Predict for test set:

- ▶ MSE for pruned: 1.922
- ▶ MSE for full: 1.945

So, pruning helped us avoid some overfitting.

(Typical pred error of $\sqrt{1.922} \approx 1.386$)

(Bigger than IQR of 1.183, but range covers $[6.98, 16.3]$.)

Example: Regression Tree

Predict for test set:

- ▶ MSE for pruned: 1.922
- ▶ MSE for full: 1.945

So, pruning helped us avoid some overfitting.

(Typical pred error of $\sqrt{1.922} \approx 1.386$)

(Bigger than IQR of 1.183, but range covers $[6.98, 16.3]$.)

(Pretty good for 1 split!?)

Random Forests

Next: random forest algorithm

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

- ▶ Boosting: models build on prior models \leadsto pick feature, predict, upweight mispredicted data, Do several times and combine.

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

- ▶ Boosting: models build on prior models \leadsto pick feature, predict, upweight mispredicted data, Do several times and combine.
- ▶ Bagging: (random select units, model) \rightarrow many times. No building.

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

- ▶ Boosting: models build on prior models \leadsto pick feature, predict, upweight mispredicted data, Do several times and combine.
- ▶ Bagging: (random select units, model) \rightarrow many times. No building.

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

- ▶ Boosting: models build on prior models \leadsto pick feature, predict, upweight mispredicted data, Do several times and combine.
- ▶ Bagging: (random select units, model) \rightarrow many times. No building.

Random Forests are bagging algorithms.

Random Forests

Next: random forest algorithm

Ensemble learning algorithms:

- ▶ Boosting: models build on prior models \leadsto pick feature, predict, upweight mispredicted data, Do several times and combine.
- ▶ Bagging: (random select units, model) \rightarrow many times. No building.

Random Forests are bagging algorithms.

Bagging: bootstrap aggregation

Random Forests

Why bag?

Random Forests

Why bag?

- ▶ Trees are low bias, high variance
(diff answers, depend on data split)

Random Forests

Why bag?

- ▶ Trees are low bias, high variance
(diff answers, depend on data split)
- ▶ Bagging averages over data subsets, reducing variance

Random Forests

Why bag?

- ▶ Trees are low bias, high variance
(diff answers, depend on data split)
- ▶ Bagging averages over data subsets, reducing variance
- ▶ (Linear regression: lower variance)

Random Forests

Random forests: decorrelated, bagged trees

Random Forests

Random forests: decorrelated, bagged trees

- ▶ Take bootstrapped training subsample

Random Forests

Random forests: decorrelated, bagged trees

- ▶ Take bootstrapped training subsample
- ▶ Build deep tree. At each split, *randomly sample m* of p predictors, build split from only those m .

Random Forests

Random forests: decorrelated, bagged trees

- ▶ Take bootstrapped training subsample
- ▶ Build deep tree. At each split, *randomly sample* m of p predictors, build split from only those m .
- ▶ (Often choose $m \approx \sqrt{p}$)

Random Forests

Random forests: decorrelated, bagged trees

- ▶ Take bootstrapped training subsample
- ▶ Build deep tree. At each split, *randomly sample* m of p predictors, build split from only those m .
- ▶ (Often choose $m \approx \sqrt{p}$)
- ▶ So, different splits consider different predictors

Random Forests

Random forests: decorrelated, bagged trees

- ▶ Take bootstrapped training subsample
- ▶ Build deep tree. At each split, *randomly sample* m of p predictors, build split from only those m .
- ▶ (Often choose $m \approx \sqrt{p}$)
- ▶ So, different splits consider different predictors
- ▶ So, trees will look very different to each other

Example: Random Forests

```
library(randomForest)

# Full bag:
bag_mps <- randomForest(ln.net ~ ., data = mp_train,
                        ntree = 500, mtry = 4,
                        importance = TRUE)

# Decorrelate:
rf_mps <- randomForest(ln.net ~ ., data = mp_train,
                      ntree = 500, mtry = 2,
                      importance = TRUE)
```

Example: Random Forests

Predict:

```
preds_bag <- predict(bag_mps, newdata = mp_test)
preds_rf  <- predict(rf_mps, newdata = mp_test)
```

- ▶ MSE for RF: 1.995
- ▶ MSE for full bag: 2.536

Example: Random Forests

Predict:

```
preds_bag <- predict(bag_mps, newdata = mp_test)
preds_rf  <- predict(rf_mps, newdata = mp_test)
```

- ▶ MSE for RF: 1.995
- ▶ MSE for full bag: 2.536

So, decorrelating helped us avoid some overfitting to each bootstrap subsample (and thus, reduced variance).

Example: Random Forests

Predict:

```
preds_bag <- predict(bag_mps, newdata = mp_test)
preds_rf  <- predict(rf_mps, newdata = mp_test)
```

- ▶ MSE for RF: 1.995
- ▶ MSE for full bag: 2.536

So, decorrelating helped us avoid some overfitting to each bootstrap subsample (and thus, reduced variance).

(Typical pred error of $\sqrt{1.995} \approx 1.412$)

Example: Random Forests

Predict:

```
preds_bag <- predict(bag_mps, newdata = mp_test)
preds_rf  <- predict(rf_mps, newdata = mp_test)
```

- ▶ MSE for RF: 1.995
- ▶ MSE for full bag: 2.536

So, decorrelating helped us avoid some overfitting to each bootstrap subsample (and thus, reduced variance).

(Typical pred error of $\sqrt{1.995} \approx 1.412$)

(Bigger than IQR of 1.183, but range covers [6.98, 16.3].)

Heterogeneous Treatment Effects

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest
- ▶ Different effects for different groups

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest
- ▶ Different effects for different groups
 - ▶ Subgroup variability (research)

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest
- ▶ Different effects for different groups
 - ▶ Subgroup variability (research)
 - ▶ Targeting resources (campaigns, marketing)

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest
- ▶ Different effects for different groups
 - ▶ Subgroup variability (research)
 - ▶ Targeting resources (campaigns, marketing)
 - ▶ Constituency effects (public policy)

Homogeneous and Heterogeneous Effects

- ▶ Most causal inference starts at *average treatment effects*
- ▶ Average may be interesting on its own, ...but often masks assumption of *homogeneous* effects
- ▶ Notationally, often assume $\tau_i = \tau \quad \forall i$
- ▶ But, *heterogeneous* effects often of central interest
- ▶ Different effects for different groups
 - ▶ Subgroup variability (research)
 - ▶ Targeting resources (campaigns, marketing)
 - ▶ Constituency effects (public policy)
- ▶ Notationally, $\exists i : \tau_i \neq \tau$

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \epsilon$$

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \epsilon$$

```
lm_out <- lm(ln.net ~ is_winner, data = mps)
lm_out
```

Call:

```
lm(formula = ln.net ~ is_winner, data = mps)
```

Coefficients:

(Intercept)	is_winner
12.2464	0.5176

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

```
t.test(ln.net ~ is_winner, data = mps)
```

Welch Two Sample t-test

data: ln.net by is_winner

t = -3.9552, df = 287.65, p-value = 9.636e-05

alternative hypothesis: true difference in means between

95 percent confidence interval:

-0.7751044 -0.2599998

sample estimates:

mean in group 0 mean in group 1

12.24641

12.76396

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \sum \beta_j X_j + \epsilon$$

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \sum \beta_j X_j + \epsilon$$

```
lm_out <- lm(ln.net ~ is_winner + is_labour +  
             is_london + age, data = mps)  
lm_out
```

Call:

```
lm(formula = ln.net ~ is_winner + is_labour + is_london + age,  
    data = mps)
```

Coefficients:

(Intercept)	is_winner	is_labour	is_london	age
12.078838	0.398818	-0.477549	0.161134	0.000000

Homogeneous and Heterogeneous Effects: Estimation

Homogeneous effects:

```
lm_lin(ln.net ~ is_winner, covariates = ~ is_labour + is_london)
```

	Estimate	Std. Error	t value
(Intercept)	1.226687e+01	0.078894901	155.4836617
is_winner	3.459885e-01	0.131207672	2.6369536
is_labour_c	-1.613663e-01	0.152608515	-1.0573871
is_london_c	2.427360e-01	0.250214401	0.9701118
age_c	4.740367e-03	0.007031323	0.6741786
is_winner:is_labour_c	-9.104022e-01	0.264395760	-3.4433313
is_winner:is_london_c	-8.847770e-02	0.426241818	-0.2075763
is_winner:age_c	-4.778657e-05	0.012753800	-0.0037468

	CI Lower	CI Upper	DF
(Intercept)	12.111785723	12.42195044	416
is_winner	0.088075873	0.60390123	416
is_labour_c	-0.461346226	0.13861367	416
is_london_c	-0.249106208	0.73457813	416

age_c	0.000000000	0.01256172	416
-------	-------------	------------	-----

Homogeneous and Heterogeneous Effects: Detection

Heterogeneous effects:

CATEs: Conditional ATEs

- ▶ *Conditional average treatment effect* (CATE):
avg treatment effect for subset of population

CATEs: Conditional ATEs

- ▶ *Conditional average treatment effect* (CATE):
avg treatment effect for subset of population
- ▶ Sometimes “CACE”

CATEs: Conditional ATEs

- ▶ *Conditional average treatment effect* (CATE):
avg treatment effect for subset of population
- ▶ Sometimes “CACE”
- ▶ Inference: not “evidence against $TE = 0$?”,
but “evidence against $CATE_1 = CATE_2$?”

Homogeneous and Heterogeneous Effects: Estimation

Heterogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \beta_2 \text{Group} + \beta_3 \text{Treatment} \cdot \text{Group} + \epsilon$$

Homogeneous and Heterogeneous Effects: Estimation

Heterogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \beta_2 \text{Group} + \beta_3 \text{Treatment} \cdot \text{Group} + \epsilon$$

► β_1 gives TE for `Group == 0`

Homogeneous and Heterogeneous Effects: Estimation

Heterogeneous effects:

$$\text{Outcome} = \beta_0 + \beta_1 \text{Treatment} + \beta_2 \text{Group} + \beta_3 \text{Treatment} \cdot \text{Group} + \epsilon$$

- ▶ β_1 gives TE for **Group** == 0
- ▶ $\beta_1 + \beta_3$ gives TE for **Group** == 1

Homogeneous and Heterogeneous Effects: Estimation

Heterogeneous effects:

```
lm_out <- lm(ln.net ~ is_winner * is_labour +  
             is_london + age, data = mps)  
coef(lm_out) |> round(3)
```

(Intercept)	is_winner	is_labour
11.959	0.780	-0.162
age	is_winner:is_labour	
0.005	-0.914	

Causal Forests

- ▶ Our regression trees had terminal nodes (“leaves”) that were sufficiently homogeneous for prediction.

Causal Forests

- ▶ Our regression trees had terminal nodes (“leaves”) that were sufficiently homogeneous for prediction.
- ▶ Use \hat{y}_{R_j} as pred value for obs in R_j

Causal Forests

- ▶ Our regression trees had terminal nodes (“leaves”) that were sufficiently homogeneous for prediction.
- ▶ Use \hat{y}_{R_j} as pred value for obs in R_j
- ▶ (Tory, winner, London \rightarrow 13.84)

Causal Forests

- ▶ Our regression trees had terminal nodes (“leaves”) that were sufficiently homogeneous for prediction.
- ▶ Use \hat{y}_{R_j} as pred value for obs in R_j
- ▶ (Tory, winner, London \rightarrow 13.84)
- ▶

$$\hat{y}_{R_j} = \frac{1}{|R_j|} \sum_{i \in R_j} Y_i$$

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment
- ▶ Controls in R_j provide good estimates of what treateds in R_j would have done under control

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment
- ▶ Controls in R_j provide good estimates of what treateds in R_j would have done under control
- ▶ Assignment w/in leaf R_j is as-good-as-random

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment
- ▶ Controls in R_j provide good estimates of what treateds in R_j would have done under control
- ▶ Assignment w/in leaf R_j is as-good-as-random
- ▶ I.e., each leaf contains an experiment

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment
- ▶ Controls in R_j provide good estimates of what treateds in R_j would have done under control
- ▶ Assignment w/in leaf R_j is as-good-as-random
- ▶ I.e., each leaf contains an experiment

Causal Forests

- ▶ Similarly, consider each leaf R_j small, homogeneous enough that potential outcomes independent
- ▶ Treateds in R_j provide good estimates of what controls in R_j would have done under treatment
- ▶ Controls in R_j provide good estimates of what treateds in R_j would have done under control
- ▶ Assignment w/in leaf R_j is as-good-as-random
- ▶ I.e., each leaf contains an experiment

$$Y(0), Y(1) \perp\!\!\!\perp T | \mathbf{X}$$

Causal Forests

► Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)
- ▶ Natural estimation of

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)
- ▶ Natural estimation of

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)
- ▶ Natural estimation of

$$\hat{\tau}_{R_j} = \frac{1}{|\{T, R_j\}|} \sum_{\{T, R_j\}} Y_i - \frac{1}{|\{C, R_j\}|} \sum_{\{C, R_j\}} Y_i$$

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)
- ▶ Natural estimation of

$$\hat{\tau}_{R_j} = \frac{1}{|\{T, R_j\}|} \sum_{\{T, R_j\}} Y_i - \frac{1}{|\{C, R_j\}|} \sum_{\{C, R_j\}} Y_i$$

So, we can use RF methods to estimate conditional (heterogeneous) treatment effects, CATEs.

Causal Forests

- ▶ Let $\{T, R_j\} = \{i : T_i = 1, i \in R_j\}$ (Tr obs in R_j)
- ▶ Let $\{C, R_j\} = \{i : T_i = 0, i \in R_j\}$ (Co obs in R_j)
- ▶ Natural estimation of

$$\hat{\tau}_{R_j} = \frac{1}{|\{T, R_j\}|} \sum_{\{T, R_j\}} Y_i - \frac{1}{|\{C, R_j\}|} \sum_{\{C, R_j\}} Y_i$$

So, we can use RF methods to estimate conditional (heterogeneous) treatment effects, CATEs.



Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees
 - ▶ Split training data into \mathcal{I} and \mathcal{J}

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees
 - ▶ Split training data into \mathcal{I} and \mathcal{J}
 - ▶ Splits chosen to maximise variance on $\hat{\tau}$ for $i \in \mathcal{J}$

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees
 - ▶ Split training data into \mathcal{I} and \mathcal{J}
 - ▶ Splits chosen to maximise variance on $\hat{\tau}$ for $i \in \mathcal{J}$
 - ▶ Splitting cannot use y_i from \mathcal{J}

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees
 - ▶ Split training data into \mathcal{I} and \mathcal{J}
 - ▶ Splits chosen to maximise variance on $\hat{\tau}$ for $i \in \mathcal{J}$
 - ▶ Splitting cannot use y_i from \mathcal{I}
 - ▶ Prediction, estimation of $\hat{\tau}$ uses only \mathcal{J}

Causal Forests: Honesty

- ▶ But need one more thing for asymptotics to work out ...
- ▶ Each tree must be *honest*
- ▶ Each unit i *either*
 - ▶ used to determine tree splits, *or*
 - ▶ used to estimate $\hat{\tau}_{R_j}$
- ▶ But **not both!**
- ▶ One way: “Double-sample” causal trees
 - ▶ Split training data into \mathcal{I} and \mathcal{J}
 - ▶ Splits chosen to maximise variance on $\hat{\tau}$ for $i \in \mathcal{J}$
 - ▶ Splitting cannot use y_i from \mathcal{J}
 - ▶ Prediction, estimation of $\hat{\tau}$ uses only \mathcal{J}
- ▶ Build a random forest (decorrelated deep trees picking from m predictors) of causal trees

Example: Causal Forests

```
library(grf)

X <- mp_train |> select(age, is_labour, is_london)

W <- mp_train |> select(is_winner) |>
  unlist() |> as.numeric()

Y <- mp_train |> select(ln.net) |> unlist()

cf_out <- causal_forest(X, Y, W)
```

Example: Causal Forests


```
X_test <- mp_test |> select(age, is_labour, is_london)

cf_pred_est_var <- predict(cf_out, X_test,
                           estimate.variance = TRUE)


cf_preds <- cf_pred_est_var$predictions

df_cf <- tibble(X_test,
                 cf_te = cf_preds,
                 cf_se = sqrt(cf_pred_est_var$variance),
                 te_lse_lower = cf_te - cf_se,
                 te_lse_upper = cf_te + cf_se)
```

Example: Causal Forests Results, Party




Example: Causal Forests Results, Party




► Mean CF TE, Tory: 0.58

Example: Causal Forests Results, Party




► Mean CF TE, Tory: 0.58 \leadsto £192,000

Example: Causal Forests Results, Party



- ▶ Mean CF TE, Tory: 0.58 \leadsto £192,000
- ▶ Mean CF TE, Labour: 0.219


Example: Causal Forests Results, Party




► Mean CF TE, Tory: 0.58 \leadsto £192,000

► Mean CF TE, Labour: 0.219 \leadsto £60,000


Example: Causal Forests Results, London



Example: Causal Forests Results, Age



Example: Causal Forests Results, Age



Variable Selection

Slide Title

Material.

Thanks!

rtm@american.edu
www.ryantmoore.org

References I

- Breiman, Leo. 2001. “Statistical Modeling: The Two Cultures.” *Statistical Science* 16 (3): 199–215. <http://www.jstor.org/stable/2676681>.
- D’Agostino McGowan, Lucy. 2023. *quartets: Datasets to Help Teach Statistics*. <https://r-causal.github.io/quartets/>.
- Hernán, Miguel A. 2018. “The c-Word: Scientific Euphemisms Do Not Improve Causal Inference from Observational Data.” *American Journal of Public Health* 108 (5): 616–19.
- Hernán, Miguel A., John Hsu, and Brian Healy. 2019. “A Second Chance to Get Causal Inference Right: A Classification of Data Science Tasks.” *CHANCE* 32 (1): 42–49. <https://doi.org/10.1080/09332480.2019.1579578>.
- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2021. *An Introduction to Statistical Learning with Applications in R*. 2nd ed. New York, NY: Springer.